

# CS410: Parallel Computing

Spring 2024

Nikhil Hegde  
Milind Chabbi

Shared-Memory Programming (using POSIX Threads)

# Recall: Why Threads?

---

- Portable, widely-available programming model
  - use on both serial and parallel systems
- Useful for hiding latency
  - e.g. latency due to I/O, communication
- Useful for scheduling and load balancing
  - especially for dynamic concurrency
- Relatively easy to program
  - significantly easier than message-passing!

# Recall:

## Shared Address Space Programming Models

---

### A brief taxonomy

- Lightweight processes and threads
  - all memory is global
  - examples: Pthreads, Cilk (lazy, lightweight threads)
- Process-based models
  - each process's data is private, unless otherwise specified
  - example: Linux shget/shmat/shmdt
- Directive-based models, e.g., OpenMP
  - shared and private data
  - facilitate thread creation and synchronization
- Global address space programming languages
  - shared and private data
  - hardware typically distributed memory, perhaps not shared
  - examples: Unified Parallel C, Co-array Fortran

# Programming Parallel Computers - Approaches

- ➡ 1. Extend existing languages:
  - Add parallel constructs (`cilk_spawn`, `omp_for` etc.)  
(requires development of compiler system support)
  - Add parallel operations (e.g. `fork`, `pthread_create`, etc.)
- ➡ 2. Define totally new parallel language and compiler system
- 3. Extend existing compilers so that they translate sequential programs into parallel programs
- 4. Add a parallel language layer on top of the sequential language

# Extend Language

- Add operations to a sequential language

- Create and terminate processes/threads
- Synchronize processes/threads
- Allow processes/threads to communicate

## **Advantages:**

- Easiest, quickest, and least expensive
- Allows existing compiler technology to be leveraged
- New libraries can be ready soon after new parallel computers are available

## **Disadvantages**

- Lack of compiler support to catch errors
- Easy to write programs that are difficult to debug

# Why Pthreads?

- Lightweight

Platform	fork()			pthread_create()		
	real	user	sys	real	user	sys
Intel 2.6 GHz Xeon E5-2670 (16 cores/node)	8.1	0.1	2.9	0.9	0.2	0.3
Intel 2.8 GHz Xeon 5660 (12 cores/node)	4.4	0.4	4.3	0.7	0.2	0.5
AMD 2.3 GHz Opteron (16 cores/node)	12.5	1.0	12.5	1.2	0.2	1.3
AMD 2.4 GHz Opteron (8 cores/node)	17.6	2.2	15.7	1.4	0.3	1.3
IBM 4.0 GHz POWER6 (8 cpus/node)	9.5	0.6	8.8	1.6	0.1	0.4
IBM 1.9 GHz POWER5 p5-575 (8 cpus/node)	64.2	30.7	27.6	1.7	0.6	1.1
IBM 1.5 GHz POWER4 (8 cpus/node)	104.5	48.6	47.2	2.1	1.0	1.5
INTEL 2.4 GHz Xeon (2 cpus/node)	54.9	1.5	20.8	1.6	0.7	0.9
INTEL 1.4 GHz Itanium2 (4 cpus/node)	54.5	1.1	22.2	2.0	1.2	0.6

[Source: Why Pthreads? | LLNL HPC Tutorials](#)

Numbers are in seconds and are obtained with 50,000 process/thread creation

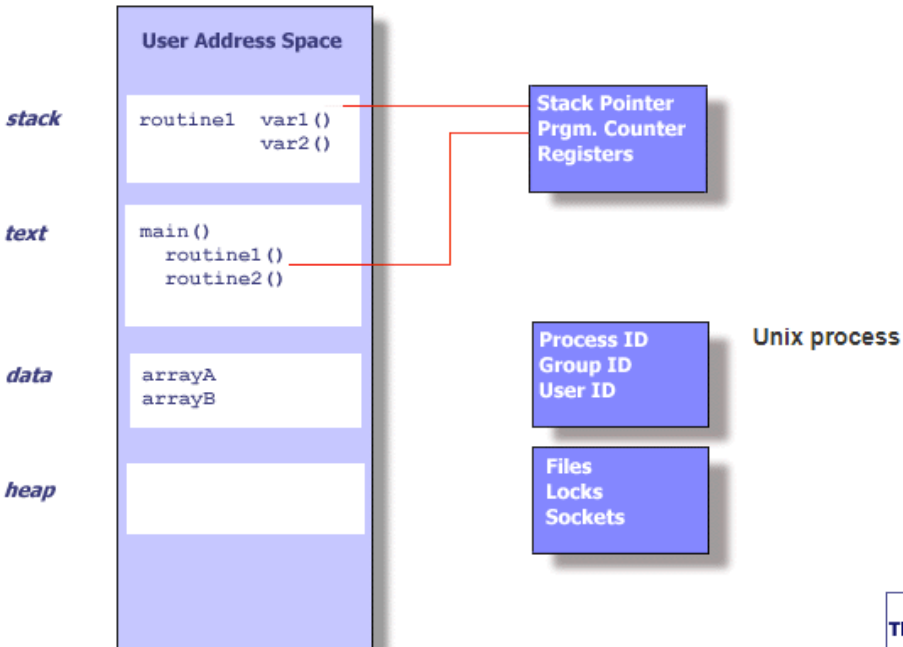
# Recap: Threads and Processes

- Abstraction provided by the OS

Process	Thread
Self-contained i.e. has its own private resources to execute/run programs. E.g. of a resource: memory.	Belongs to a process. Share memory and other resources among threads of the same process.
Is an instance of a running program.	Can be considered as a subroutine in the 'main' program
Have an illusion that <i>entire computer</i> is for itself.	Have an illusion that <i>entire processor</i> is for itself.

# Recap: Threads and Processes

[Source: what is a Thread? | LLNL HPC Tutorials](#)

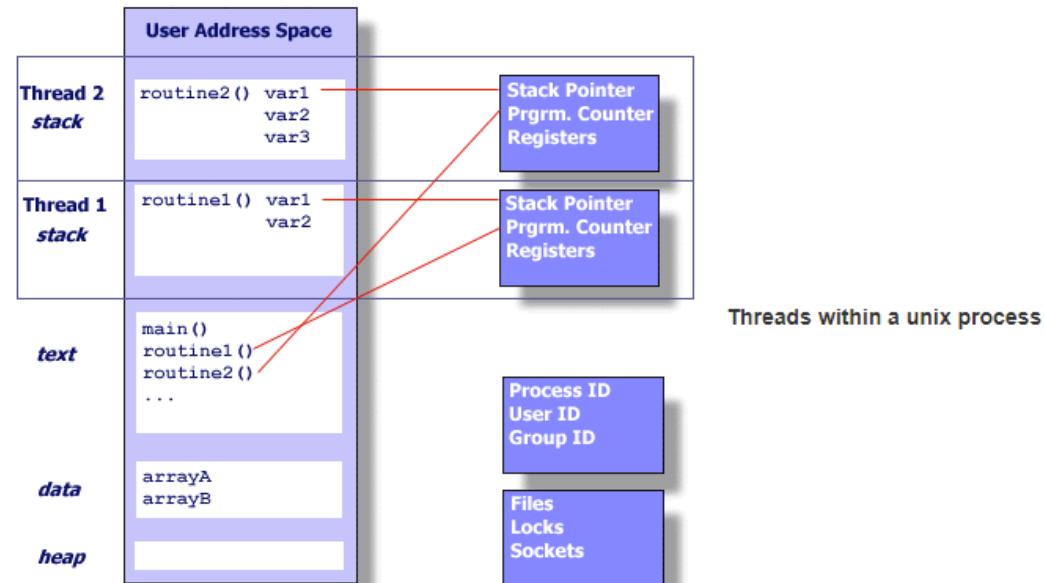


Each process maintains its own:

- Process ID, Group ID, User ID,
- Environment,
- Working directory,
- Registers, stack, heap, program instructions
- File handles, Shared libraries, ..

Each thread in a process maintains its own copy of:

- Stack pointer,
- Program counter,
- Registers,
- Thread-specific data,
- Scheduling policy





# POSIX Threads

- Hardware vendors implemented proprietary versions of threads earlier
  - Software portability was a concern
- IEEE POSIX 1003.1c standard
  - Standardized C language threads programming interface for UNIX like systems - **Pthreads API**
- Implemented as a set of C language programming types and procedure calls
  - Include a header file `pthread.h` and use a library (linking may be implicit in some cases)
  - Fortran Programmers can use wrappers around C procedure calls. Some Fortran compilers provide Fortran Pthreads API.

# POSIX Threads

- Pthreads implementation now supported by most vendors
  - In addition to proprietary implementation
- Concepts are broadly applicable to other implementations (independent of Pthreads API)
  - NT Threads
  - Solaris Threads
  - Java Threads
  - C++ Threads (C++11 onwards)

# Pthreads API

- Subroutines grouped into four major categories:

- 1. Thread Management –**

Routines that deal with thread creation, join, detach, etc. Also, query thread attributes

- 2. Mutexes –**

Routines that deal with synchronization. Creating, destroying, locking, and unlocking a mutex (abbr. for “mutual exclusion”). Also, set mutex attributes

- 3. Condition Variables –**

Routines addressing communication between threads that share a mutex. Based upon programmer specified conditions. Includes functions to create, destroy, wait and signal based upon specified variable values. Also, set/query condition variable attributes.

- 4. Synchronization**

Routines that manage read-write locks and barriers

Contains around 100 routines

# Pthreads API

- Subroutines grouped into four major categories:



## **1. Thread Management –**

Routines that deal with thread creation, join, detach, etc. Also, query thread attributes

## **2. Mutexes –**

Routines that deal with synchronization. Creating, destroying, locking, and unlocking a mutex (abbr. for “mutual exclusion”). Also, set mutex attributes

## **3. Condition Variables –**

Routines addressing communication between threads that share a mutex. Based upon programmer specified conditions. Includes functions to create, destroy, wait and signal based upon specified variable values. Also, set/query condition variable attributes.

## **4. Synchronization**

Routines that manage read-write locks and barriers

# Creating Pthreads

`pthread_create(threadID, attr, start_routine, arg)`

unique thread ID  
Returned upon  
successful creation.

Thread attributes specified  
with object of type  
`pthread_attr_t`. Can  
leave it as NULL to set  
default attribute values.

The routine that the  
thread executes once  
created.

Single argument  
passed to  
`start_routine`  
(type-casted to void  
\*). Can leave it as  
NULL to omit  
arguments.

**Creates a new thread and makes it ready-to-execute**

- Usage: Typically, `main` function first creates threads, which in turn may create other threads. No hierarchy exists among threads; All threads created are peers.
- Return value: an integer indicating status of creation

# Pthread ID `pthread_t`

```
int pthread_create(pthread_t* threadID, ..
```

- `pthread_t` object is initialized by the `pthread_create` API.

# Pthreads Attributes `pthread_attr_t`

```
int pthread_create(pthread_t* threadID,  
const pthread_attr_t* attribute,...
```

- `pthread_attr_init` and `pthread_attr_destroy` APIs create and destroy thread attribute object.

The attribute object contains options to specify:

- Detached or joinable state
- Scheduling inheritance
- Scheduling policy
- Scheduling priority
- Scheduling contention scope
- Stack size
- Stack address

# Start Routine

```
int pthread_create(pthread_t* threadID,  
const pthread_attr_t* attribute,  
void * (*start_routine)(void*),..
```

- Called upon thread execution



# Arguments to Start Routine

```
int pthread_create(pthread_t* threadID,  
const pthread_attr_t* attribute,  
void * (*start_routine)(void*),  
void * arg)
```

- Single argument typecasted to void \*
- How do you pass multiple arguments?
  - Pack the arguments in a structure object and pass the pointer to the structure object (cast as void \*)
- Since thread start time is non-deterministic, do not pass as arguments data that is modified by other threads

# Terminating Pthreads

- `pthread_exit(status)`
- `pthread_cancel(thread)` routine can be used by a thread to cancel another thread
- `exit(int)` terminates the entire process (including all threads)

# pthread\_exit(status)

- `pthread_exit()` routine is called after a thread has completed its work and is no longer required to exist
- If `main()` finishes before the threads it has created, and exits with `pthread_exit()`, the other threads will continue to execute.
  - Otherwise, they will be automatically terminated when `main()` finishes
- The programmer may optionally specify a termination *status*, which is stored as a void pointer for any thread that may join the calling thread
- Cleanup
  - `pthread_exit()` routine does not close files
  - Recommended to use `pthread_exit()` to exit from all threads...especially `main()`.

# Joining and Detaching from Pthreads

- `pthread_join(threadid, status)` is one way to accomplish synchronization between threads
  - Blocks the calling thread until the specified thread terminates
  - Only threads created as joinable can be joined (other option is to create a thread as *detached*, where the thread can never be joined.)
    - Can Use `pthread_detach()` to detach a thread after it was created as joinable

