

CS410: Parallel Computing

Spring 2024

Nikhil Hegde
Milind Chabbi

Parallel Algorithms and Applications

Embarrassingly Parallel Applications


- Application where:
 - A number of (almost) independent tasks
 - No or very little communication between tasks
 - Each task can be executed on a node
- Master-worker approach could be used
- Examples
 - Image Processing: e.g. blurring, scaling, rotation etc.
 - Computer Graphics: e.g. ray tracing
 - Monte Carlo method: e.g. estimation of pi
 - ...

Interesting reads:

<http://graphics.pixar.com/library/CurlyHairA/paper.pdf>,

<https://www.fxguide.com/fxfeatured/brave-new-hair/>

Image Processing



Pixel

- 8-bits – 256 colors possible.
- 24-bits – More than 16 million colors possible
- Voxel – 3 dimensional image

Simple Image Processing

- Scaling


- Scale the image by a factor λ_x in x – direction

$$x' = x \cdot \lambda_x$$

$$y' = y$$

- Scaling Matrix R: $\begin{bmatrix} \lambda_x & 0 \\ 0 & 1 \end{bmatrix}$

usually 3D: $\begin{bmatrix} \lambda_x & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$ and $\begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$ (*3rd dim is a constant, usually 1*)



https://en.wikipedia.org/wiki/Digital_image_processing

- The computation is done for all pixels.
 - Notice that computation at each pixel does not depend on any other data other than the pixel value

Simple Image Processing

- Shifting an object


$$x' = x + \Delta x \quad y' = y + \Delta y$$

- Rotation of an object

$$x' = x \cos\theta - y \sin\theta \quad y' = x \sin\theta + y \cos\theta$$


Rotation Matrix R:
$$\begin{bmatrix} \cos\theta & -\sin\theta & 0 \\ \sin\theta & \cos\theta & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = R \times \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$



https://en.wikipedia.org/wiki/Digital_image_processing


Simple Image Processing



$$x(i, j) = f\{ x(i, j), x(i, j + 1), x(i, j - 1), x(i + 1, j), \\ x(i - 1, j) \}$$

Example: Average values of the neighbouring pixels

Simple Image Processing



$$P'(i,j) = f \{ P(i,j), P(i+1, j+1), P(i+1, j), P(i+1, j-1), P(i-1, j), P(i-1, j-1), P(i-1, j), P(i-1, j+1), P(i, j+1) \}$$

$P'(i,j)$ – new value of the pixel (i,j)



Simple Image Processing

- Lowpass (e.g. usage: blurring)

Mask / Kernel R: $\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$

$$P'_{ij} = 1/9 \{ P_{ij} + P_{i+1j} + P_{i-1j} + P_{ij-1} + P_{i-1j-1} + P_{i+1j-1} + P_{ij+1} + P_{i-1j+1} + P_{i+1j+1} \}$$

P'_{ij} – new value of the pixel (i,j)



https://en.wikipedia.org/wiki/Digital_image_processing

Highpass Kernel: Mask / Kernel R: $\frac{1}{9} \begin{bmatrix} 0 & -1 & 0 \\ -1 & 4 & -1 \\ 0 & -1 & 0 \end{bmatrix}$


Simple Image Processing: Parallelizing

- Static Task Assignment
 - Divide the image region into fixed number of partitions
 - Assign each region to a distinct process / thread / processor
 - Different pixels may require different amount computation / iterations (depending on the application)
 - Number of iterations required – not known *a priori*
 - Unbalanced load for different processes
 - Performance – not very good

Dynamic Task Assignment

- Dynamic (i.e. runtime) allocation of tasks work pool – computations of different regions


Simple Image Processing: Parallelizing




Assume: one Pixel Processing per Processor and neighbour pixel data is required in an application considered

Inter-processor Communication Requirements ?


Data Decomposition



Square regions



Horizontal strips






Vertical strips

Assume: region-wise partitioning of the pixel grid and neighbour pixel data is required in the application considered

Inter-processor Communication Requirements ?

Data Decomposition & Inter-processor Communication

- 4 Processors



Computations/processor

\propto (Area of the partition)


Communications/processor

\propto (\sum length of the partitions adjacent to other partitions)

Monte Carlo Methods

- Computations based on pseudo-random or quasi-random numbers
Example: For C programs, use `rand()` for generating integer pseudo-random numbers; first use `srand()` to seed the `rand()` function.
- Applications: radiation transport, Monte Carlo simulation in communications, solution of partial differential equations (PDEs)
- Each process (processor) requires a pseudo-random (or quasi-random) number generator
- Obtain an estimate of the solution
- Estimate converges to the solution as the number of trials are increased

Monte Carlo Method



Circle

radius $r = 1$

Area of circle = $\pi r^2 = \pi$

Area of square = $2 \times 2 = 4$

$$\frac{\text{Area of circle}}{\text{Area of square}} = \frac{\pi}{4}$$


$$\pi = 4 \cdot (\text{area of circle}) / (\text{area of square})$$

Monte Carlo Method: Estimate of π

- Choose points randomly within the square
(x,y): choose x and y coordinates randomly for a point
- Estimate of $\pi = \frac{4(\text{Number of points in circle})}{(\text{Number of points in square})}$
- Estimate of $\pi \rightarrow$ true value of π as
number points $\rightarrow \infty$
- Computations for all points are independent \rightarrow can be done in
parallel
- How to generate pseudo-random numbers in parallel?
- OpenMP Program (Demo)

Application: Laplace Equation

- Second order PDE : $\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} = 0$



- Example
 - Metal plate of size 6 cm x 10 cm
 - Each edge (boundary) is held at a constant temperature
 - Find temperatures of points within the plate
 - Steady-state solution

Application: Laplace Equation

- We begin by writing difference equation for approximating the PDE
- Discretize the region (create a mesh of grid points)
- Compute the temperature at each grid point

Laplace Equation – Numerical Solution

1. Approximate the derivatives of $\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} = 0$ using central differences

$$\frac{\partial^2 u}{\partial x^2} \approx \frac{(u(x + \delta x, y) - 2u(x, y) + u(x - \delta x, y))}{(\delta x)^2}$$

$$\frac{\partial^2 u}{\partial y^2} \approx \frac{(u(x, y + \delta y) - 2u(x, y) + u(x, y - \delta y))}{(\delta y)^2}$$

Where, δx and δy are step sizes along x and y direction resp.

Laplace Equation – Numerical Solution

- Substituting in $\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} = 0$:

$$\frac{(u(x + \delta x, y) - 2u(x, y) + u(x - \delta x, y))}{(\delta x)^2}$$

+

$$\frac{(u(x, y + \delta y) - 2u(x, y) + u(x, y - \delta y))}{(\delta y)^2}$$


=

$$\frac{(u(x + \delta x, y) + u(x, y + \delta y) - 4u(x, y) + u(x - \delta x, y) + u(x, y - \delta y))}{(h)^2}$$

$$= 0$$

Laplace Equation – Numerical Solution


- Representing $u(x, y)$



Notation: $u_{i,j}$

Laplace Equation – Numerical Solution


- Representing $u(x - \delta x, y)$



Notation: $u_{i-1,j}$

Laplace Equation – Numerical Solution


- Representing $u(x + \delta x, y)$



Notation: $\mathbf{u}_{i+1,j}$

Laplace Equation – Numerical Solution


- Representing $u(x, y - \delta y)$



Notation: $\mathbf{u}_{i,j-1}$

Laplace Equation – Numerical Solution

- Representing $u(x, y + \delta y)$



Notation: $u_{i,j+1}$


Laplace Equation – Numerical Solution

- Rewriting using notation:

$$\frac{(u(x + \delta x, y) + u(x, y + \delta y) - 4u(x, y) + u(x - \delta x, y) + u(x, y - \delta y))}{(h)^2}$$

$$= 0$$

$$\frac{u_{i+1,j} + u_{i,j+1} - 4u_{i,j} + u_{i-1,j} + u_{i,j-1}}{h^2} = 0$$



called 5-point stencil



Laplace Equation – Numerical Solution

- Consider the *boundary-value* problem:

$$\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} = 0 \text{ in the square } 0 < x < 1, 0 < y < 1$$

$$u = x^2 y \text{ on the boundary, } h = 1/3$$

$$\frac{u_{i+1,j} + u_{i,j+1} - 4u_{i,j} + u_{i-1,j} + u_{i,j-1}}{h^2} = 0 \quad \text{Equation 1}$$




Laplace Equation – Numerical Solution

- Substituting for i, j in equation 1 and computing u_{11}

$$(u_{i+1,j} + u_{i,j+1} - 4u_{i,j} + u_{i-1,j} + u_{i,j-1} = 0)$$

$$u_{21} + u_{12} - 4u_{11} + u_{01} + u_{10} = 0$$

$$u_{21} + u_{12} - 4u_{11} + 0 + 0 = 0$$




Laplace Equation – Numerical Solution

- Computing u_{21}

$$(u_{i+1,j} + u_{i,j+1} - 4u_{i,j} + u_{i-1,j} + u_{i,j-1} = 0)$$

$$u_{31} + u_{22} - 4u_{21} + u_{11} + u_{20} = 0$$

$$\frac{1}{3} + u_{22} - 4u_{21} + u_{11} + 0 = 0$$




Laplace Equation – Numerical Solution

- Computing u_{12}

$$(u_{i+1,j} + u_{i,j+1} - 4u_{i,j} + u_{i-1,j} + u_{i,j-1} = 0)$$

$$u_{22} + u_{13} - 4u_{12} + u_{02} + u_{11} = 0$$

$$u_{22} + 1/9 - 4u_{12} + 0 + u_{11} = 0$$




Laplace Equation – Numerical Solution

- Computing u_{22}

$$(u_{i+1,j} + u_{i,j+1} - 4u_{i,j} + u_{i-1,j} + u_{i,j-1} = 0)$$

$$u_{32} + u_{23} - 4u_{22} + u_{12} + u_{21} = 0$$

$$\frac{2}{3} + \frac{4}{9} - 4u_{22} + u_{12} + u_{21} = 0$$



Laplace Equation – Numerical Solution

• System of Equations

$$(u_{i+1,j} + u_{i,j+1} - 4u_{i,j} + u_{i-1,j} + u_{i,j-1} = 0)$$

Right

Top

Center

Left


Bottom

$$u_{21} + u_{12} - 4u_{11} + 0 + 0 = 0$$

$$1/3 + u_{22} - 4u_{21} + u_{11} + 0 = 0$$

$$u_{22} + 1/9 - 4u_{12} + 0 + u_{11} = 0$$

$$2/3 + 4/9 - 4u_{22} + u_{12} + u_{21} = 0$$



Laplace Equation – Numerical Solution

- Computing System of Equations:

$$u_{21} + u_{12} - 4u_{11} + 0 + 0 = 0$$

$$1/3 + u_{22} - 4u_{21} + u_{11} + 0 = 0$$

$$u_{22} + 1/9 - 4u_{12} + 0 + u_{11} = 0$$

$$2/3 + 4/9 - 4u_{22} + u_{12} + u_{21} = 0$$

$$\begin{pmatrix} -4 & 1 & 1 & 0 \\ 1 & -4 & 0 & 1 \\ 1 & 0 & -4 & 1 \\ 0 & 1 & 1 & -4 \end{pmatrix} \begin{pmatrix} u_{11} \\ u_{21} \\ u_{12} \\ u_{22} \end{pmatrix} = \begin{pmatrix} 0 \\ -1/3 \\ -1/9 \\ -10/9 \end{pmatrix} \quad \mathbf{Ax=B}$$

A
x
=
B

Matrix A has only coefficients

$$\begin{pmatrix} 1 & 1 & 1 & 1 \\ 1 & -4 & 1 & 1 \\ 1 & 1 & 1 & 1 \end{pmatrix}$$

Computing Stencil – Iterative Methods

- Jacobi and Gauss-Seidel
 - Start with an initial guess for the unknowns u^0_{ij}
 - Improve the guess u^1_{ij}
 - Iterate: derive the new guess, u^{n+1}_{ij} , from old guess u^n_{ij}

Background – Jacobi Iteration

- **Goal:** find solution to system of equations represented by $AX=B$
- **Approach:** find sequence of approximations $X^0, X^1, X^2, \dots, X^n$, which gradually approach X .
 - X^0 is called initial guess, X^i 's called *iterates*
- **Method:**
 - Split A into $A=L+D+U$ e.g.

$$\begin{pmatrix} -4 & 1 & 1 & 0 \\ 1 & -4 & 0 & 1 \\ 1 & 0 & -4 & 1 \\ 0 & 1 & 1 & -4 \end{pmatrix} = \begin{pmatrix} 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 \end{pmatrix} + \begin{pmatrix} -4 & 0 & 0 & 0 \\ 0 & -4 & 0 & 0 \\ 0 & 0 & -4 & 0 \\ 0 & 0 & 0 & -4 \end{pmatrix} + \begin{pmatrix} 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{pmatrix}$$

\uparrow \uparrow \uparrow
L D U

Background – Jacobi Iteration

- **Compute:** $AX=B$ is $(L+D+U)X=B$

$$\Rightarrow DX = -(L+U)X+B$$

$$\Rightarrow DX^{(k+1)} = -(L+U)X^k+B \quad \textbf{(iterate step)}$$

$$\Rightarrow X^{(k+1)} = D^{-1} (-(L+U)X^k) + D^{-1}B$$

(As long as D has no zeros in the diagonal $X^{(k+1)}$ is obtained)

- E.g.
$$\begin{pmatrix} -4 & 0 & 0 & 0 \\ 0 & -4 & 0 & 0 \\ 0 & 0 & -4 & 0 \\ 0 & 0 & 0 & -4 \end{pmatrix} \begin{pmatrix} u_{11}^{\mathbf{1}} \\ u_{21} \\ u_{12} \\ u_{22} \end{pmatrix} = - \begin{pmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{pmatrix} \begin{pmatrix} u_{11}^{\mathbf{0}} \\ u_{21} \\ u_{12} \\ u_{22} \end{pmatrix} + \begin{pmatrix} 0 \\ -1/3 \\ -1/9 \\ -10/9 \end{pmatrix},$$


u_{ij} 's value in ($\mathbf{1}$)st iteration is computed based on u_{ij} values computed in ($\mathbf{0}$)th iteration

Background – Jacobi Iteration

- E.g.
$$\begin{pmatrix} -4 & 0 & 0 & 0 \\ 0 & -4 & 0 & 0 \\ 0 & 0 & -4 & 0 \\ 0 & 0 & 0 & -4 \end{pmatrix} \begin{pmatrix} u_{11} \\ u_{21} \\ u_{12} \\ u_{22} \end{pmatrix}^{k+1} = - \begin{pmatrix} 0 & 1 & 1 & 0 \\ 1 & 0 & 0 & 1 \\ 1 & 0 & 0 & 1 \\ 0 & 1 & 1 & 0 \end{pmatrix} \begin{pmatrix} u_{11} \\ u_{21} \\ u_{12} \\ u_{22} \end{pmatrix}^k + \begin{pmatrix} 0 \\ -1/3 \\ -1/9 \\ -10/9 \end{pmatrix},$$

u_{ij} 's value in $(k+1)^{st}$ iteration is computed based on u_{ij} values computed in $(k)^{th}$ iteration

- Center's value is updated. Why?



5-point stencil

Computing Stencil

- Jacobi Solution Approach:
 - Approximate the *value of the center* with *old values* of (left, right, top, bottom)


Computing Stencil

- $u_{right} + u_{top} - 4u_{center} + u_{left} + u_{bottom} = 0$
 $\Rightarrow u_{center} = 1/4(u_{right} + u_{top} + u_{left} + u_{bottom})$
- Applying Jacobi Iteration:
$$u_{center}^{(k+1)} = 1/4(u_{right}^{(k)} + u_{top}^{(k)} + u_{left}^{(k)} + u_{bottom}^{(k)})$$

Computing Stencil

- Example: applying Jacobi Iteration for 6x6 grid:

$$u_{center}^{(k+1)} = 1/4(u_{right}^{(k)} + u_{top}^{(k)} + u_{left}^{(k)} + u_{bottom}^{(k)})$$




Iteration 1

1) Compute u_{11} using initial guess for u_{12} and u_{21} . u_{01} and u_{10} are known from boundary conditions

Computing Stencil

- Example: applying Jacobi Iteration:

$$u_{center}^{(k+1)} = 1/4(u_{right}^{(k)} + u_{top}^{(k)} + u_{left}^{(k)} + u_{bottom}^{(k)})$$



Iteration 1


1) Compute u_{11} using initial guess for u_{12} and u_{21} . u_{01} and u_{10} are known from boundary conditions

2) Compute u_{21} using initial guess for u_{11} , u_{31} , and u_{22} . u_{20} are known from boundary conditions


In 2), note that the initial guess for u_{11} is used even though u_{11} was updated just before in 1)

Computing Stencil


- In every iteration, suppose we follow the computing order as shown (dashed):




In any iteration, what are all the points of a 5-point stencil already updated while computing u_{ij} ?



Computing Stencil



What are the points that are already computed at $u_{i,j}$?



Parallel Implementation

- Using new values as soon as they are computed
- Cannot be parallelized
- Solution
 - Use all old values to find new values for all internal mesh points in parallel
 - Red-black ordering
 - Wavefront ordering

Parallel Implementation

- Assuming 1 grid point per process and Jacobi (solution)
- Each process sends and receives 4 messages (except the processes adjacent to a boundary)

Total number of iterations = m

Total number of messages per process = $m (4 + 4)$
= $8 m$ send/receive messages

Total number of messages for a corner process = $4 m$

Total number of messages for a process adjacent to only one boundary point = $6 m$

Laplace's Equation: Finite Difference Method

– Red-Black Ordering

✕ - Red ● - Black

✕	●	✕	●	✕	●
---	---	---	---	---	---

●	✕	●	✕	●	✕
---	---	---	---	---	---

✕	●	✕	●	✕	●
---	---	---	---	---	---

●	✕	●	✕	●	✕
---	---	---	---	---	---

Step1: Compute Red points

Step2: Compute Black points

Step3: If convergence not reached, go back to step1

Laplace's Equation: Monte Carlo Method

- Discretize the region, mesh size $h = \frac{1}{4}$
- Random walk
 - Start at the point where solution is desired
 - Generate a random number between $(0,1)$ with uniform distribution



- Decide the direction of the next step based on the value of random number
- Continue the random walk until arrive at a boundary point. Note that boundary value b_i for i -th random walk

Laplace's Equation: Monte Carlo Method

- Carry out N random walks, until the estimate of solution sufficiently converges

$$\text{Estimate of the solution} = \frac{1}{N} \sum_{i=1}^N b_i$$

- Estimate \rightarrow solution, as $N \rightarrow \infty$

Monte Carlo Method: Laplace's Equation

- Each random walk is independent of the other random walks
 - All random walks can be carried out in parallel
 - Each process (processor) would require a pseudo-random number generator
- The estimate of the solution converges to the approximate solution based on the grid size (h)
- The solution can be obtained only at a particular point inside the region
- Number of steps in a random walk- random variable
- Dynamic task allocation preferable when number of processors $< N$

Laplace's Equation: Monte Carlo Solution

Reference:

Bhavsar, V.C. and J.R. Isaac, "Design and Analysis of Parallel Monte Carlo Algorithms"

SIAM J. Statistical and Scientific Computing,
Vol. 8, No. 1, s73-s95, Jan. 1987.